

Kubernetes Networking: Load Balancer vs. Ingress

Introduction

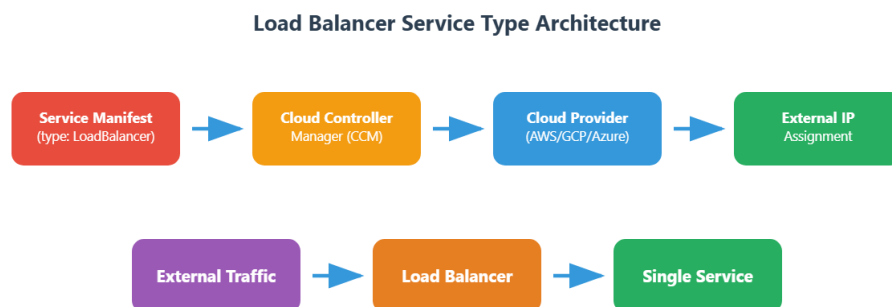
Kubernetes is a powerful platform for managing containerized applications, and effective networking is crucial for exposing services to external users. In my recent studies through Abhishek Veermala's DevOps course on Udemy, I explored two key methods for managing external traffic: **load balancer service type** and **Ingress**. Understanding their differences is essential for designing scalable, cost-effective, and flexible systems, which are critical skills for DevOps and cloud-native roles.

Load Balancer Service Type

How It Works

The **load balancer service type** is a Kubernetes service configuration that exposes a service to external traffic by provisioning a cloud provider's load balancer (e.g., AWS ALB). Here's the process:

- A service manifest is created with `type: LoadBalancer`.
- The Kubernetes **API server** passes the request to the **Cloud Controller Manager (CCM)**.
- The CCM communicates with the cloud provider (e.g., AWS) to create a load balancer.
- An **external IP** is assigned, allowing external access to the service.



Advantages

- **Simplicity:** Requires only a single manifest change (`type: LoadBalancer`), reducing setup complexity.

Disadvantages

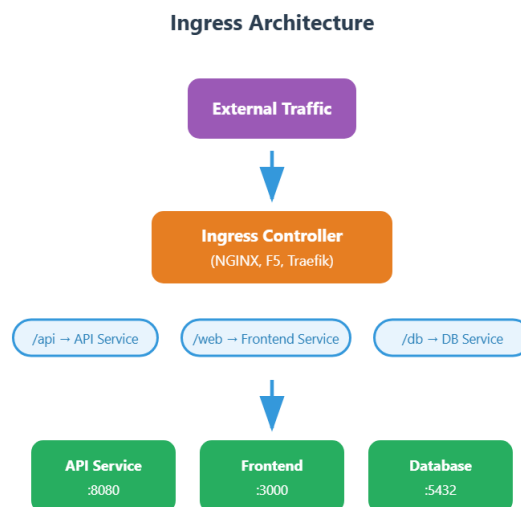
- **Non-declarative Configuration:** Changes like switching to HTTPS or adjusting load balancer settings must be made manually in the cloud provider's console, not in Kubernetes manifests. This lacks version control and traceability.
- **Cost Ineffectiveness:** Each service with `type: LoadBalancer` gets its own load balancer, which can be expensive for multiple services (e.g., 10 services = 10 load balancers).
- **Limited Flexibility:** Tied to the cloud provider's load balancer (e.g., AWS ALB), with no support for alternatives like NGINX or F5.
- **Dependency on CCM:** Only works on clusters with a CCM, failing on local setups like Minikube.

Ingress

How It Works

Ingress is a Kubernetes API object that manages external access to services, typically HTTP/HTTPS, using a single entry point. It requires an **ingress controller** (e.g., NGINX, F5) to handle traffic routing based on rules defined in the Ingress resource. For example:

- External traffic hits the ingress controller.
- The controller routes traffic to services based on host or path rules (e.g., `/api` to an API service).



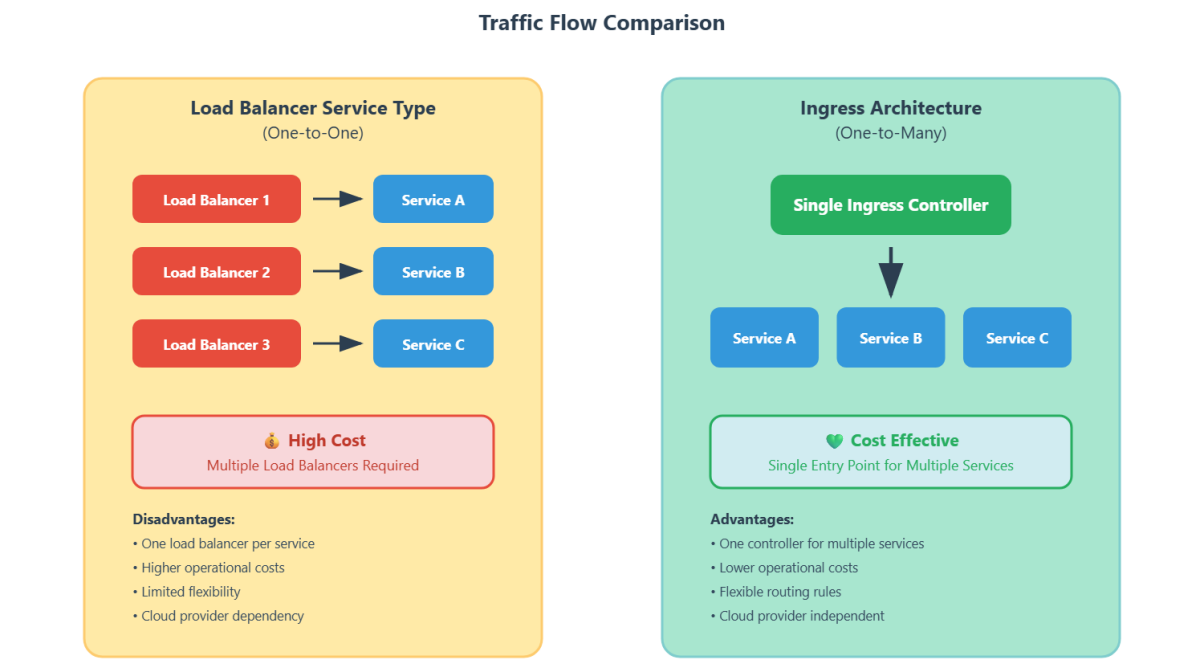
Advantages

- **Declarative Configuration:** Defined via YAML files, enabling version control and reproducibility.
- **Cost-Effective:** A single ingress can route traffic to multiple services using path-based or host-based routing, reducing the need for multiple load balancers.
- **Flexible:** Supports various ingress controllers (e.g., NGINX, F5, Traefik), allowing customization.
- **Universal Compatibility:** Works on any Kubernetes cluster, including local setups like Minikube, without requiring a CCM.

Configuration Example

Below is a sample Ingress YAML:

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: ingress-example
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx-example
  rules:
  - http:
      paths:
      - path: /api
        pathType: Prefix
        backend:
          service:
            name: api-service
            port:
              number: 80
```



Comparison Table

Aspect	Load Balancer Service Type	Ingress
Configuration	Non-declarative, manual changes	Declarative via YAML, version-controlled
Cost	One load balancer per service, costly	One ingress for multiple services, cost-effective

Aspect	Load Balancer Service Type	Ingress
Flexibility	Tied to CCM (e.g., AWS ALB)	Supports various ingress controllers (e.g., NGINX, F5)
Compatibility	Requires CCM, doesn't work on Minikube	Works on any cluster, including Minikube
Ease of Use	Simple to configure	Requires additional setup (ingress YAML, controller)

Conclusion

While the **load balancer service type** is simple to configure, its limitations—cost, lack of flexibility, and dependency on cloud infrastructure—make it less suitable for complex, scalable applications. **Ingress**, with its declarative nature, cost efficiency, and flexibility, is the preferred choice for production environments. These insights are critical for DevOps engineers aiming to optimize Kubernetes deployments.