

KUBERNETES DAEMONSETS, JOBS & CRONJOBS

DAEMONSETS:

1. Daemonsets ensures that exactly one pod is running on each node of the cluster.
2. If New node is added to the cluster, Daemonset will create a new pod on that node automatically.
3. If an existing node is removed from the cluster, Daemonset will delete its pod from that node.
4. If we try to delete a deamonset pod manually, Daemonset controller creates a new pod instantly on that node.

Use Cases for DaemonSets:

DaemonSets are typically used to deploy **cluster-wide system services** such as:

- Monitoring agents (e.g., Prometheus Node Exporter, Datadog Agent)
- Logging agents (e.g., Fluentd, Fluent Bit, Filebeat)
- Networking plugins (e.g., CNI plugins like AWS VPC CNI, Azure CNI, Calico, Flannel, Cilium)
- Storage plugins (e.g., CSI Node Drivers like AWS EBS CSI, AWS EFS CSI, GCP PD CSI etc)
- Security agents (e.g., Falco runtime security monitors)
- Telemetry collectors (e.g., OpenTelemetry Collector DaemonSet)

Example of Daemonset:

kube-proxy, a critical system component responsible for **Service-to-Pod networking** inside the cluster.

```
kubectl get ds -n kube-system
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
kindnet	3	3	3	3	3	kubernetes.io/os=linux	10m
kube-proxy	3	3	3	3	3	kubernetes.io/os=linux	10m

Kube-proxy is running as daemonset in the K8s cluster.

```
kubectl get pods -n kube-system -o wide | grep -i proxy
```

```
root@DESKTOP-C6P8EQS:/kubernetes/18)Daemonset_Job_CronJ$ kubectl get pods -n kube-system -o wide | grep -i proxy
kube-proxy-59rcp           1/1   Running   0      47m  172.18.0.4  rayeez-cluster-control-plane  <none>       <none>
kube-proxy-c8wrf           1/1   Running   0      46m  172.18.0.3  rayeez-cluster-worker    <none>       <none>
kube-proxy-z9lgx           1/1   Running   0      46m  172.18.0.2  rayeez-cluster-worker2  <none>       <none>
```

But how it was possible to deploy daemonset pod of kube-proxy on the control plane node?

Control plain node is tainted by default as follows in **KIND** cluster:

```
kubectl describe nodes rayeez-cluster-control-plane
```

```
Name:                 rayeez-cluster-control-plane
Roles:               control-plane
Labels:              beta.kubernetes.io/arch=amd64
                     beta.kubernetes.io/os=linux
                     kubernetes.io/arch=amd64
                     kubernetes.io/hostname=rayeez-cluster-control-plane
                     kubernetes.io/os=linux
                     node-role.kubernetes.io/control-plane=
                     node.kubernetes.io/exclude-from-external-load-balancers=
Annotations:         kubeadm.alpha.kubernetes.io/cri-socket: unix:///run/containerd/containerd.sock
                     node.alpha.kubernetes.io/ttl: 0
                     volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp:   Fri, 05 Dec 2025 09:10:39 +0000
Taints:              node-role.kubernetes.io/control-plane:NoSchedule
Unschedulable:        false
Lease:
  HolderIdentity:   rayeez-cluster-control-plane
  AcquireTime:      <unset>
  RenewTime:        Fri, 05 Dec 2025 10:10:07 +0000
```

Control plain node is tainted but even though daemonset pod of kube-proxy is scheduled on it because it has following tolerations:

```
kubectl get ds -n kube-system kube-proxy -o yaml
```

```
  serviceAccountName: kube-proxy
  terminationGracePeriodSeconds: 30
  tolerations:
  - operator: Exists
  volumes:
  - configMap:
```

This means kube-proxy daemonset pod can tolerate any taint, which make it possible to deployed on control plain node.

Note: System-level DaemonSets are typically deployed into their own dedicated namespace for better organization and access control.

Kube-proxy daemonsets pods are deployed in kube-system namespace:

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  annotations:
    deprecated.daemonset.template.generation: "1"
  creationTimestamp: "2025-12-05T09:10:43Z"
  generation: 1
  labels:
    k8s-app: kube-proxy
  name: kube-proxy
  namespace: kube-system
  resourceVersion: "509"
  uid: 62fbe917-77c3-4c76-89f0-ce447f5b69a5
```

Lets Explore the Concept of Daemonset practically!

ds.yaml:

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: log-collector
  namespace: logging-ns
  labels:
    app: log-collector
spec:
  selector:
    matchLabels:
      app: log-collector
  template:
    metadata:
      labels:
        app: log-collector
  spec:
    tolerations:
      - key: "node-role.kubernetes.io/control-plane"
        operator: "Exists"
        effect: "NoSchedule"
    containers:
      - name: log-collector
```

```

image: busybox
command:
  [
    "/bin/sh",
    "-c",
    "while true; do echo 'Collecting
logs...'; sleep 30; done",
  ]
resources:
  requests:
    cpu: "50m"
    memory: "50Mi"
  limits:
    cpu: "100m"
    memory: "100Mi"
volumeMounts:
  - name: varlog
    mountPath: /var/log
volumes:
  - name: varlog
    hostPath:
      path: /var/log
      type: Directory

```

Description of above ds.yaml manifest:

- A Daemonset named log-collector will be deployed in logging-ns namespace.
- This will create and manage pods having label log-collector.
- Daemonset pod is having tolerations matching with Control plane node taint. Hence daemonset pod will also scheduled on control plain node.
- Container inside the pod will run a infinite while loop printing 'Collecting logs... ' every 30 seconds.

- CPU, Memory resources requests and limits are also defined for the container to provide proper resources to the daemonset pods and prevent it starving other pods.
- A HostPath volume is also defined and mounted on the container, mapping the node's /var/log directory to the container's /var/log directory to simulate a logging agent's behavior.

As system level Daemonset have their own dedicated namespace, First we have to create the namespace for it:

```
kubectl create ns logging-ns
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get ns
NAME          STATUS  AGE
default       Active  31m
kube-node-lease Active  31m
kube-public   Active  31m
kube-system   Active  31m
local-path-storage Active  30m
logging-ns    Active  6s
```

Lets apply the ds.yaml file:

```
kubectl apply -f ds.yaml
```

Verify;

```
kubectl get ds -n logging-ns
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get ds -n logging-ns
NAME      DESIRED  CURRENT  READY  UP-TO-DATE  AVAILABLE  NODE SELECTOR  AGE
log-collector  3        3        0      3           0          <none>   14s
```

Let make logging-ns namespace as default for the current context:

```
kubectl config set-context --current --namespace=logging-ns
```

Now I can list out daemonsets for logging-ns namespace directly:

```
kubectl get ds
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get ds
NAME      DESIRED  CURRENT  READY  UP-TO-DATE  AVAILABLE  NODE SELECTOR  AGE
log-collector  3        3        3      3           3          <none>   4m26s
```

```
kubectl get pods
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
log-collector-5h9km  1/1     Running   0          45s
log-collector-652pb  1/1     Running   0          45s
log-collector-gmxxpm 1/1     Running   0          45s
```

These are the Daemonset pods of log-collector in logging-ns namespace.

```
kubectl describe ds log-collector
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl describe ds log-collector
Name:           log-collector
Namespace:      logging-ns
Selector:       app=log-collector
Node-Selector:  <none>
Labels:         app=log-collector
Annotations:    deprecated.daemonset.template.generation: 1
Desired Number of Nodes Scheduled: 3
Current Number of Nodes Scheduled: 3
Number of Nodes Scheduled with Up-to-date Pods: 3
Number of Nodes Scheduled with Available Pods: 3
Number of Nodes Misscheduled: 0
Pods Status:   3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=log-collector
  Containers:
    log-collector:
      Image:      busybox
      Port:       <none>
      Host Port: <none>
      Command:
        /bin/sh
        -c
        while true; do echo 'Collecting logs...'; sleep 30; done
      Limits:
        cpu:        100m
        memory:    100Mi
      Requests:
        cpu:        50m
        memory:    50Mi
      Environment: <none>
      Mounts:
        /var/log from varlog (rw)
  Volumes:
    varlog:
      Type:      HostPath (bare host directory volume)
      Path:      /var/log
      HostPathType: Directory
    Node-Selectors: <none>
  Tolerations:  node-role.kubernetes.io/control-plane:NoSchedule op=Exists
Events:
  Type  Reason          Age   From            Message
  ----  ----          --   --              --
  Normal SuccessfulCreate 10m   daemonset-controller  Created pod: log-collector-wxwzg
  Normal SuccessfulCreate 10m   daemonset-controller  Created pod: log-collector-kjzhl
  Normal SuccessfulCreate 10m   daemonset-controller  Created pod: log-collector-nrb9b
```

Lets delete one of the daemonset pod and watch it from another terminal:

```
kubectl delete pods log-collector-5h9km
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl delete pods log-collector-5h9km
pod "log-collector-5h9km" deleted from logging-ns namespace
```

Watch;

```
kubectl get pods -w
```

NAME	READY	STATUS	RESTARTS	AGE
log-collector-5h9km	1/1	Running	0	60s
log-collector-652pb	1/1	Running	0	60s
log-collector-gmxpm	1/1	Running	0	60s
log-collector-5h9km	1/1	Terminating	0	86s
log-collector-5h9km	0/1	Error	0	115s
log-collector-kgl86	0/1	Pending	0	0s
log-collector-kgl86	0/1	Pending	0	0s
log-collector-kgl86	0/1	ContainerCreating	0	0s
log-collector-5h9km	0/1	Error	0	116s
log-collector-5h9km	0/1	Error	0	116s
log-collector-kgl86	1/1	Running	0	11s

Termination of this existing daemonset pod and Creation of new daemonset pod takes place simultaneously. This is performed by Daemonset controller.

JOB:

1. Job is a kubernetes object that runs one-off tasks such as copying a file from location A to location B, ping an IP address such as 8.8.8.8 to check connectivity, Deleting temporary files etc.
2. Once the task finishes successfully (Pod status = Completed), the Job is considered finished.
3. It creates one or more Pods to run a specific task to completion.
4. Jobs are useful for:
 - Running batch jobs, Database migrations, clearing unused files or creating backups reliably and One-time report generation.

Lets Explore Job practically!

job.yaml:

```
apiVersion: batch/v1
kind: Job
metadata:
  name: hello-job
spec:
```

```
ttlSecondsAfterFinished: 60
completions: 2
parallelism: 2
backoffLimit: 4
template:
  spec:
    containers:
      - name: hello
        image: busybox
        command: ["bin/sh", "-c", "echo Hello
from the Job! && sleep 10"]
    restartPolicy: Never
```

ttlSecondsAfterFinished: 60

This means delete the job 60 seconds after the task is finished(Either successfully or Failed)

completions: 2

This means Job runs 2 pods for the task to be succeed. If 1 Pod has successfully completed the and second pod failed the task, Then Job status is marked as failed.

parallelism: 2

This means run 2 pods concurrently, If parallelism is 1, it means run pods sequentially one after another.

backoffLimit: 4

This means If task is failing, recreate the pod till 4 attempts, If it still fails, Job is considered as failed. This recreation of pods is performed by **Job Controller**.

kubelet will restart the container within the same Pod on the same node depending on **restartPolicy** specified.

Container inside each pods runs echo command and sleep for 10 seconds.

Apply the job.yaml manifest:

```
kubectl apply -f job.yaml  
kubectl get job -w
```

NAME	STATUS	COMPLETIONS	DURATION	AGE
hello-job	Running	0/2		0s
hello-job	Running	0/2	0s	0s
hello-job	Running	0/2	6s	6s
hello-job	Running	0/2	16s	16s
hello-job	Running	0/2	17s	17s
hello-job	Complete	2/2	17s	17s
hello-job	Complete	2/2	17s	77s
hello-job	Complete	2/2	17s	77s

Once the Job is completed successfully, It removed after 60 seconds.

```
kubectl get pods -w
```


NAME	READY	STATUS	RESTARTS	AGE
hello-job-v69dp	0/1	Pending	0	0s
hello-job-f6n2b	0/1	Pending	0	0s
hello-job-v69dp	0/1	Pending	0	0s
hello-job-f6n2b	0/1	Pending	0	0s
hello-job-v69dp	0/1	ContainerCreating	0	0s
hello-job-f6n2b	0/1	ContainerCreating	0	0s
hello-job-f6n2b	1/1	Running	0	4s
hello-job-v69dp	1/1	Running	0	6s
hello-job-f6n2b	0/1	Completed	0	13s
hello-job-f6n2b	0/1	Completed	0	14s
hello-job-v69dp	0/1	Completed	0	15s
hello-job-f6n2b	0/1	Completed	0	15s
hello-job-v69dp	0/1	Completed	0	16s
hello-job-v69dp	0/1	Completed	0	16s
hello-job-v69dp	0/1	Completed	0	76s
hello-job-f6n2b	0/1	Completed	0	76s
hello-job-v69dp	0/1	Completed	0	76s
hello-job-f6n2b	0/1	Completed	0	76s

Pods status converted to from running to completed in 10 seconds and then after 60 seconds Job is removed from the cluster.

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get job  
NAME      STATUS  COMPLETIONS DURATION  AGE  
hello-job  Complete 2/2          17s       76s  
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get job  
No resources found in default namespace.
```

CRONJOB:

1. Jobs are scheduled to run on specified time using Cronjob.



2. At each scheduled time, the CronJob controller automatically creates a new Job object, and that Job is responsible for running one or more Pods to complete the task.
3. Once created, the Job behaves like any other Kubernetes Job — handling retries, backoffLimit, and managing success or failure.

Lets explore the concept of Cronjob practically!

cronjob.yaml:

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: hello-cronjob
spec:
  schedule: "*/1 * * * *"    # Cron expression
  specifying the job to run every 1 minute
  jobTemplate:
    spec:
      backoffLimit: 4
      ttlSecondsAfterFinished: 60
      completions: 2
      parallelism: 2
      template:
        spec:
          containers:
            - name: hello
              image: busybox
```

```

        command: ["bin/sh", "-c", "echo
Hello from CronJob! && sleep 10"]
      restartPolicy: Never

```

Lets Apply this yaml;

```
kubectl apply -f cronjob.yaml
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get cj
NAME          SCHEDULE      TIMEZONE    SUSPEND   ACTIVE   LAST SCHEDULE   AGE
hello-cronjob  */1 * * * *  <none>     False      0        23s           3m57s
```

```
kubectl get jobs
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get jobs
No resources found in default namespace.
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get jobs
NAME          STATUS      COMPLETIONS   DURATION   AGE
hello-cronjob-29417979  Running    0/2           6s         6s
```

Jobs is scheduled to run every minute.

```
kubectl get pods
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get pods
NAME          READY   STATUS      RESTARTS   AGE
hello-cronjob-29417979-cn7cs  0/1     Completed   0          63s
hello-cronjob-29417979-j4r2b  0/1     Completed   0          63s
hello-cronjob-29417980-4brxk  0/1     ContainerCreating   0          3s
hello-cronjob-29417980-zh6bz  0/1     ContainerCreating   0          3s
```

New pods are created after a minute as per the schedule.

```
root@DESKTOP-C6P8EQS:~/kubernetes/18)Daemonset_Job_CronJ$ kubectl get pods
NAME          READY   STATUS      RESTARTS   AGE
hello-cronjob-29417980-4brxk  0/1     Completed   0          27s
hello-cronjob-29417980-zh6bz  0/1     Completed   0          27s
```

Previous Job is removed as per the TTL seconds timer and pods are garbage collected by K8s.

The CronJob creates a new Job every minute. Once each Job completes—whether it succeeds or fails—it's automatically cleaned up based on the TTL set in the manifest.

Without a TTL value, the CronJob keeps creating Jobs every minute, causing completed Pods to accumulate in the cluster and clutter the environment.

```

Ubuntu                               Ubuntu
root@DESKTOP-C6P8EQS:/kubernetes$ kubectl get jobs -w
root@DESKTOP-C6P8EQS:/kubernetes$ kubectl get pods -w
NAME      STATUS  COMPLETIONS DURATION   AGE
hello-cronjob-29417994  Running  0/2        14s       14s
hello-cronjob-29417994  Running  0/2        15s       15s
hello-cronjob-29417994  Running  0/2        16s       16s
hello-cronjob-29417994  Complete 2/2        16s       16s
hello-cronjob-29417995  Running  0/2        0s        0s
hello-cronjob-29417995  Running  0/2        0s        0s
hello-cronjob-29417995  Running  0/2        4s        4s
hello-cronjob-29417995  Running  0/2        7s        7s
hello-cronjob-29417995  Running  0/2        14s       14s
hello-cronjob-29417995  Running  0/2        15s       15s
hello-cronjob-29417995  Running  1/2        16s       16s
hello-cronjob-29417995  Running  1/2        19s       19s
hello-cronjob-29417995  Complete 2/2        19s       19s
hello-cronjob-29417996  Running  0/2        0s        0s
hello-cronjob-29417996  Running  0/2        0s        0s
hello-cronjob-29417996  Running  0/2        5s        5s
hello-cronjob-29417996  Running  0/2        7s        7s
hello-cronjob-29417996  Running  0/2        15s       15s
hello-cronjob-29417996  Running  0/2        16s       16s
hello-cronjob-29417996  Running  1/2        16s       16s
hello-cronjob-29417996  Running  1/2        17s       17s
hello-cronjob-29417996  Running  1/2        18s       18s
hello-cronjob-29417996  Complete 2/2        18s       18s
hello-cronjob-29417997  Running  0/2        0s        0s
hello-cronjob-29417997  Running  0/2        0s        0s
hello-cronjob-29417997  Running  0/2        6s        6s
hello-cronjob-29417997  Running  0/2        16s       16s
hello-cronjob-29417997  Running  0/2        17s       17s
hello-cronjob-29417997  Running  1/2        17s       17s
hello-cronjob-29417997  Running  1/2        18s       18s
hello-cronjob-29417997  Complete 2/2        18s       18s
hello-cronjob-29417994  Complete 2/2        3m18s    3m18s
hello-cronjob-29417998  Running  0/2        0s        0s
hello-cronjob-29417998  Running  0/2        0s        0s
hello-cronjob-29417998  Running  0/2        5s        5s
hello-cronjob-29417998  Running  0/2        6s        6s
hello-cronjob-29417998  Running  0/2        16s       16s

                                         ^Croot@DESKTOP-C6P8EQS:/kubernetes/18)Daemonset_Job_CronJ$ kubectl get pods -w
NAME      READY  STATUS  RESTARTS  AGE
hello-cronjob-29417994-hh9c8  0/1  Pending  0  0s
hello-cronjob-29417994-blk4p  0/1  Pending  0  0s
hello-cronjob-29417994-hh9c8  0/1  Pending  0  0s
hello-cronjob-29417994-blk4p  0/1  Pending  0  0s
hello-cronjob-29417994-blk4p  0/1  ContainerCreating 0  0s
hello-cronjob-29417994-blk4p  0/1  ContainerCreating 0  0s
hello-cronjob-29417994-blk4p  1/1  Running  0  4s
hello-cronjob-29417994-blk4p  1/1  Running  0  4s
hello-cronjob-29417994-blk4p  0/1  Completed 0  14s
hello-cronjob-29417994-hh9c8  0/1  Completed 0  14s
hello-cronjob-29417994-blk4p  0/1  Completed 0  15s
hello-cronjob-29417994-hh9c8  0/1  Completed 0  15s
hello-cronjob-29417994-blk4p  0/1  Completed 0  16s
hello-cronjob-29417994-blk4p  0/1  Completed 0  16s
hello-cronjob-29417995-sp8c7  0/1  Pending  0  0s
hello-cronjob-29417995-sp8c7  0/1  Pending  0  0s
hello-cronjob-29417995-t29lt  0/1  Pending  0  0s
hello-cronjob-29417995-t29lt  0/1  Pending  0  0s
hello-cronjob-29417995-sp8c7  0/1  ContainerCreating 0  0s
hello-cronjob-29417995-t29lt  0/1  ContainerCreating 0  0s
hello-cronjob-29417995-sp8c7  1/1  Running  0  3s
hello-cronjob-29417995-t29lt  1/1  Running  0  6s
hello-cronjob-29417995-sp8c7  0/1  Completed 0  13s
hello-cronjob-29417995-sp8c7  0/1  Completed 0  15s
hello-cronjob-29417995-sp8c7  0/1  Completed 0  15s
hello-cronjob-29417995-t29lt  0/1  Completed 0  16s
hello-cronjob-29417995-t29lt  0/1  Completed 0  18s
hello-cronjob-29417995-t29lt  0/1  Completed 0  19s
hello-cronjob-29417996-knrgj  0/1  Pending  0  0s
hello-cronjob-29417996-knrgj  0/1  Pending  0  0s
hello-cronjob-29417996-ntnpv  0/1  Pending  0  0s
hello-cronjob-29417996-ntnpv  0/1  Pending  0  0s
hello-cronjob-29417996-ntnpv  0/1  ContainerCreating 0  0s
hello-cronjob-29417996-ntnpv  0/1  ContainerCreating 0  0s
hello-cronjob-29417996-knrgj  1/1  Running  0  3s
hello-cronjob-29417996-knrgj  1/1  Running  0  6s
hello-cronjob-29417996-knrgj  0/1  Completed 0  14s

```