


## REQUESTS, LIMITS & LIMIT-RANGES



Everything runs smoothly until deployment-2 malfunctions. A code bug causes its pods to consume excessive CPU and memory, starving deployment-1 resources and creating a noisy-neighbor problem.

This is where the concept of requests and limits acts as a safeguard.


**Request:** Minimum guaranteed resources (typically CPU & Memory) for a container.

**Limits:** Maximum resources (typically CPU & Memory) a container can use as long as the resources are available from the node.

### Benefits of Using Requests and Limits:

1. Efficient Resource Allocation
2. Avoidance of resource Starvation - Prevents one workload from hogging all the resources.
3. Cluster Stability
4. Mitigation of Noisy Neighbour problem - A container that suddenly consumes excessive node resources becomes a noisy neighbor, and limits help prevent it from impacting other workloads.

**Note:** It is recommended to define Requests and Limits for each container within a pod.



- If a node has 6 CPUs and 24 GB of memory, and six pods deployed, each request 1 CPU and 4 GB, the node becomes fully allocated. A seventh pod cannot be scheduled because no remaining resources are available.
- Pods can use the amount of resources they request, and if they need more, they can consume up to their defined limit—provided the node has available capacity.

### What happens when a container goes beyond resource limits?

**CPU:** The **KERNEL** THROTTLES the CPU.

**Memory:** The **KERNEL** may Terminate the pod.

- **Memory Pressure (no memory available):** The pod is **OOM-killed**.
- **No Memory Pressure:** The pod continues using memory beyond its limit.

- Lets Explore the concept of Request and Limits Practically!


First, we need to set up the metrics server on our cluster. We can do this by downloading the required manifest from the Kubernetes GitHub repository:

`"kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml"`

And we have to add this flag to the deployment manifest in this file:

```
--kubelet-insecure-tls
```

```
spec:  
  containers:  
    - args:  
        - --cert-dir=/tmp  
        - --secure-port=10250  
        - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname  
        - --kubelet-use-node-status-port  
        - --metric-resolution=15s  
        - --kubelet-insecure-tls  
      image: registry.k8s.io/metrics-server/metrics-server:v0.8.0
```



- The Kubelet binary includes a component called **cAdvisor** (Container Advisor), which exposes an API that collects metrics from nodes and pods and provides them to the Metrics Server.

Lets deploy metric server on the cluster:

```
kubectl apply -f components.yaml
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl apply -f components.yaml  
serviceaccount/metrics-server created  
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created  
clusterrole.rbac.authorization.k8s.io/system:metrics-server created  
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created  
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created  
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created  
service/metrics-server created  
deployment.apps/metrics-server created  
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
```

## **kubectl get pods -n kube-system**

```
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl get pods -n kube-system
NAME                               READY   STATUS    RESTARTS   AGE
coredns-7c65d6fcf9-89rxf          1/1    Running   0          55m
coredns-7c65d6fcf9-w4px7          1/1    Running   0          55m
etcd-rayeez-cluster-control-plane  1/1    Running   0          56m
kindnet-c6zxx                     1/1    Running   0          55m
kindnet-kx97f                     1/1    Running   0          55m
kindnet-snjk                         1/1    Running   0          55m
kube-apiserver-rayeez-cluster-control-plane 1/1    Running   0          56m
kube-controller-manager-rayeez-cluster-control-plane 1/1    Running   0          56m
kube-proxy-ns5l4                   1/1    Running   0          55m
kube-proxy-sspcd                   1/1    Running   0          55m
kube-proxy-zljdz                   1/1    Running   0          55m
kube-scheduler-rayeez-cluster-control-plane 1/1    Running   0          56m
metrics-server-bf688598-gp85q      1/1    Running   0          3m16s
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ |
```

Lets verify resource usage by nodes and pods:

## **kubectl top nodes**

```
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl top nodes
NAME           CPU(cores)   CPU(%)   MEMORY(bytes)   MEMORY(%)
rayeez-cluster-control-plane  191m        1%       818Mi        22%
rayeez-cluster-worker         32m        0%       212Mi        5%
rayeez-cluster-worker2        26m        0%       177Mi        4%
```

## **kubectl top pods**

```
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl top pods
NAME           CPU(cores)   MEMORY(bytes)
frontend-deploy-6df94c895b-fwsps  12m        10Mi
frontend-deploy-6df94c895b-gpspl  0m         10Mi
frontend-deploy-6df94c895b-msmvx  14m        10Mi
frontend-deploy-6df94c895b-rlr4v  0m         10Mi
frontend-deploy-6df94c895b-tmkjp  12m        10Mi
```

## **kubectl describe nodes rayeez-cluster-worker2**

| Namespace   | Name                             | CPU Requests | CPU Limits | Memory Requests | Memory Limits | Age   |
|-------------|----------------------------------|--------------|------------|-----------------|---------------|-------|
| default     | frontend-deploy-6df94c895b-gpspl | 0 (0%)       | 0 (0%)     | 0 (0%)          | 0 (0%)        | 2m25s |
| default     | frontend-deploy-6df94c895b-rlr4v | 0 (0%)       | 0 (0%)     | 0 (0%)          | 0 (0%)        | 2m25s |
| default     | frontend-deploy-6df94c895b-tmkjp | 0 (0%)       | 0 (0%)     | 0 (0%)          | 0 (0%)        | 2m25s |
| kube-system | kindnet-snjk                     | 100m (0%)    | 100m (0%)  | 50Mi (1%)       | 50Mi (1%)     | 61m   |
| kube-system | kube-proxy-zljdz                 | 0 (0%)       | 0 (0%)     | 0 (0%)          | 0 (0%)        | 61m   |

Table for Resource requests and Limits can be observed:

- Let apply the following memory.yaml which simulates the memory load on the container with 190Mi:

```
apiVersion: v1
kind: Pod
metadata:
  name: memory-demo
spec:
  containers:
```

```

- name: memory-demo-ctr
  image: polinux/stress
  resources:
    requests:
      memory: "100Mi"
    limits:
      memory: "200Mi"
  command: ["stress"]
  args: [--vm, "1", "--vm-bytes", "190M",
         "--vm-hang", "1"]

```

Apply;

```
kubectl apply -f memory.yaml
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
frontend-deploy-6df94c895b-fwsp 1/1   Running   0          7m25s  10.244.2.3  rayeez-cluster-worker   <none>        <none>
frontend-deploy-6df94c895b-gpspl 1/1   Running   0          7m25s  10.244.1.4  rayeez-cluster-worker2  <none>        <none>
frontend-deploy-6df94c895b-msmvx 1/1   Running   0          7m25s  10.244.2.4  rayeez-cluster-worker   <none>        <none>
frontend-deploy-6df94c895b-rlr4v  1/1   Running   0          7m25s  10.244.1.2  rayeez-cluster-worker2  <none>        <none>
frontend-deploy-6df94c895b-tmkjp  1/1   Running   0          7m25s  10.244.1.3  rayeez-cluster-worker2  <none>        <none>
memory-demo   1/1   Running   0          18s    10.244.1.5  rayeez-cluster-worker2  <none>        <none>
```

```
kubectl top pods
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl top pods
NAME          CPU(cores)   MEMORY(bytes)
frontend-deploy-6df94c895b-fwsp  0m          10Mi
frontend-deploy-6df94c895b-gpspl 0m          10Mi
frontend-deploy-6df94c895b-msmvx 0m          10Mi
frontend-deploy-6df94c895b-rlr4v  0m          10Mi
frontend-deploy-6df94c895b-tmkjp  0m          10Mi
memory-demo   390m        190Mi
```

→ Let's change this argument from **190Mi** to **250Mi**, which exceeds the container's defined limit of **200Mi** and apply a new pod.

Apply;

```
kubectl apply -f memory.yaml
```


```
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
frontend-deploy-6df94c895b-fwsp  1/1   Running   0          14m    10.244.2.3  rayeez-cluster-worker   <none>        <none>
frontend-deploy-6df94c895b-gpspl 1/1   Running   0          14m    10.244.1.4  rayeez-cluster-worker2  <none>        <none>
frontend-deploy-6df94c895b-msmvx 1/1   Running   0          14m    10.244.2.4  rayeez-cluster-worker   <none>        <none>
frontend-deploy-6df94c895b-rlr4v  1/1   Running   0          14m    10.244.1.2  rayeez-cluster-worker2  <none>        <none>
frontend-deploy-6df94c895b-tmkjp  1/1   Running   0          14m    10.244.1.3  rayeez-cluster-worker2  <none>        <none>
memory-demo   1/1   Running   0          7m15s  10.244.1.5  rayeez-cluster-worker2  <none>        <none>
memory-demo-2  0/1   OOMKilled  0          6s    10.244.2.7  rayeez-cluster-worker   <none>        <none>
```

Pod Status is **OOMKilled**. But Nodes have enough resources available.

```
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl top nodes
NAME          CPU(cores)   CPU(%)   MEMORY(bytes)   MEMORY(%)
rayeez-cluster-control-plane  209m       1%      784Mi        21%
rayeez-cluster-worker        61m        0%      300Mi        8%
rayeez-cluster-worker2       218m      1%      456Mi        12%
```

It is clear that there is no **Memory Pressure**. Then Why pod is getting **OOMKilled**?

→ **Concept of Control Groups(A Linux Concept):**



1. Control groups essentially create resource boundaries, ensuring that each process is allowed to use only a defined amount of resources—such as memory, CPU, and I/O—without affecting others.
2. If a pod attempts to consume more resources than allowed, the **kernel** will terminate the process.
3. When the container runtime creates and starts a container inside a pod, it also sets up a corresponding **control group** for that container.
4. Each running pod corresponds to a process managed by a control group. If the **cgroup** detects that a pod is consuming more memory than what is defined in its resource specification, it terminates that process. This results in the pod entering the **OOMKilled** (Out of Memory Killed) state.

Lets create a Pod which simulates the CPU load of 1000m which beyond the limit specified(900m):

```
apiVersion: v1
kind: Pod
metadata:
  name: cpu-demo
spec:
  containers:
    - name: cpu-container
```

```

image: vish/stress
resources:
  requests:
    cpu: "500m"
  limits:
    cpu: "900m"
args:
  - --cpus
  - "1"

```

Apply;

```
kubectl apply -f cpu.yaml
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl top pods
NAME                               CPU(cores)   MEMORY(bytes)
cpu-demo                           899m        1Mi
frontend-deploy-6df94c895b-fwsp  0m          10Mi
frontend-deploy-6df94c895b-gpspl 0m          10Mi
frontend-deploy-6df94c895b-msmrx 0m          10Mi
frontend-deploy-6df94c895b-rlr4v  0m          10Mi
frontend-deploy-6df94c895b-tmkjp 0m          10Mi
memory-demo                         179m        190Mi
```

CPU Utilization is not going beyond 900m.

The **kernel throttles** CPU usage when a container attempts to consume more CPU resources than the limits defined for it.

```
kubectl describe nodes rayeez-cluster-worker2
```

| Namespace   | Name                             | CPU Requests | CPU Limits | Memory Requests | Memory Limits | Age   |
|-------------|----------------------------------|--------------|------------|-----------------|---------------|-------|
| default     | cpu-demo                         | 500m (4%)    | 900m (7%)  | 0 (0%)          | 0 (0%)        | 6m27s |
| default     | frontend-deploy-6df94c895b-gpspl | 0 (0%)       | 0 (0%)     | 0 (0%)          | 0 (0%)        | 40m   |
| default     | frontend-deploy-6df94c895b-rlr4v | 0 (0%)       | 0 (0%)     | 0 (0%)          | 0 (0%)        | 40m   |
| default     | frontend-deploy-6df94c895b-tmkjp | 0 (0%)       | 0 (0%)     | 0 (0%)          | 0 (0%)        | 40m   |
| default     | memory-demo                      | 0 (0%)       | 0 (0%)     | 100Mi (2%)      | 200Mi (5%)    | 33m   |
| kube-system | kindnet-snjk                     | 100m (0%)    | 100m (0%)  | 50Mi (1%)       | 50Mi (1%)     | 99m   |
| kube-system | kube-proxy-zljdz                 | 0 (0%)       | 0 (0%)     | 0 (0%)          | 0 (0%)        | 99m   |

Lets Explore **LimitRange** Object of K8s:

→ **LimitRange**: Setting Default Requests and Limits

A LimitRange enforces default CPU and memory requests/limits in a **namespace**, automatically applying them when containers don't specify their own.

Let apply the following LimitRange yaml:

```

apiVersion: v1
kind: LimitRange
metadata:
  name: resource-limits
  namespace: default
spec:
  limits:
    - type: Container
      default:
        cpu: "2" # Default CPU limit
        memory: "4Gi" # Default memory limit
      defaultRequest:
        cpu: "1" # Default CPU request
        memory: "2Gi" # Default memory request
      max:
        cpu: "4" # Maximum CPU a container can
        request
        memory: "8Gi" # Maximum memory a
        container can request
      min:
        cpu: "500m" # Minimum CPU a container
        must request
        memory: "512Mi" # Minimum memory a
        container must request

```

Apply;

```
kubectl apply -f limitrange.yaml
```

Verify;

```
kubectl describe limits resource-limits
```

```

root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl describe limits resource-limits
Name:      resource-limits
Namespace: default
Type       Resource  Min     Max   Default Request  Default Limit  Max Limit/Request Ratio
----       -----   ---     ---   -----          -----          -----
Container  cpu       500m    4     1             2             -
Container  memory   512Mi   8Gi   2Gi          4Gi          -

```

Lets create a pod imperatively and check its specifications:

```

kubectl run test-pod --image=nginx
kubectl describe pods test-pod

```

```

Ready:          True
Restart Count: 0
Limits:
  cpu:    2
  memory: 4Gi
Requests:
  cpu:     1
  memory: 2Gi
Environment: <none>

```

Requests and Limits are set for this pod because of LimitRange applied in default namespace

→ Lets try to create a pod less than 500m CPU:

```

apiVersion: v1
kind: Pod
metadata:
  name: cpu-demo-2
spec:
  containers:
    - name: cpu-container
      image: vish/stress
      resources:
        requests:
          cpu: "400m"
        limits:
          cpu: "900m"
      args:
        - --cpus
        - "1"

```

Apply;

```
kubectl apply -f cpu.yaml
```

```
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ kubectl apply -f cpu.yaml
Error from server (Forbidden): error when creating "cpu.yaml": pods "cpu-demo-2" is forbidden: minimum cpu usage per Container is 500m, but request is 400m
root@DESKTOP-C6P8EQS:~/kubernetes/10)Request_Limits_LimitRange$ |
```

It gives an error: minimum cpu usage per Container is **500m**, but request is **400m**.

This happens because the LimitRange enforces a minimum CPU request of 500m for every container.